



Quicker Sorting

Dietmar Kühl

Quicker Sorting

Dietmar Kühl
API Technology London
Bloomberg L.P.

Overview

- quickly sorting out Quick-Sort
- implement and compare variations of Quick-Sort
- turning the concrete implementation generic


C++ Standard Library

- complexity of C++ Standard Library sort algorithms is $O(n \log n)$
- `std::sort()` - in-place, fast
- `std::stable_sort()` - stable, not in-place
- `std::make_heap()/std::sort_heap()`

Data Set

- multiple random sequence
 - with varying degrees of equal values
- one each ascending/descending sequence
- one sequence with all the same value

Results



Quick-Sort

input



Quick-Sort

input



pivot



Quick-Sort


input



pivot



$x < p$



Quick-Sort


input




pivot



$x < p$



partition



Quick-Sort


input




pivot




$x < p$



partition



recurse



Simple Implementation

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1;  
    int* mid = std::partition(begin, pivot,  
                             [=](int value){ return value < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```

Results



Use a Pivot

- popular choice: median of 3 (or more)
- works well with sorted or reverse sorted inputs
- adds extra costs to determine the median
- ... using last is bad, though


Pivot Implementation

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1, * mid = begin + len / 2;  
    small::sort(begin, pivot, mid);  
    mid = std::partition(begin, pivot,  
                        [=](int arg){ return arg < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```

Sorting 3 Elements

```
template <typename I>
static inline void small::sort(I a, I b, I c) {
    if (*b < *a) {
        if (*c < *b) std::iter_swap(a, c);
        else if (*c < *a) { auto t(*a); *a = *b; *b = *c; *c = t; }
        else std::iter_swap(a, b);
    } else {
        if (*c < *a) { auto t(*c); *c = *b; *b = *a; *a = t; }
        else if (*c < *b) std::iter_swap(b, c);
    }
}
```


Results




Intro Sort

- monitor depth of recursion:
 - no more than $2 \log n$ recursive calls
- too deep \Rightarrow fallback to a different algorithm
 - in-place required: heap sort
 - enough spare memory: merge sort

Intro Sort

```
void sort(int* begin, int* end, int depth, int max) {  
    if (++depth < max) {  
        ...  
        sort(begin, mid, depth, max);  
        sort(mid + 1, end, depth, max);  
    }  
    else { std::stable_sort(begin, end); }  
}  
  
void sort(int* begin, int* end) {  
    int s(end-begin), m(0); while (s >>= 1) { ++m; }  
    sort(begin, end, 0, 2 * m);  
}
```

Results




Insertion Sort

- quick sort isn't really effective on small ranges
- insertion sort is more effective in that case
- two potential approaches:
 - stop when ranges too small and post-process
 - sort small ranges with insertion sort

Local Insertion Sort

```
if (len <= Size) {  
    if (begin == end) return;  
    for (int* it = begin; ++it != end; )  
        if (*it < *(it - 1)) {  
            int* c = it, t = *c;  
            *c = *(c - 1);  
            while (begin != --c && t < *(c - 1))  
                *c = *(c - 1);  
            *c = t;  
        }  
    return;  
}
```

Results




“Tail Recursion”

- using more stack is slower
- instead of call use a loop for the bigger part
- equivalent to tail call optimisation
 - however, C++ compilers are not too strong

Tail Recursion

```
while (20 < (len = std::distance(begin, end))) {  
    if (++depth < max) {  
        ...  
        if (mid - begin < end - mid) {  
            sort(begin, mid, depth, max);  
            begin = mid + 1; }  
        else {  
            sort(mid + 1, end, depth, max);  
            end = mid; }  
        }  
    }
```

Results




Partition

- so far using `std::partition()`:
`mid = std::partition(begin, end, predicate);`
- really the core part of the algorithm


Partition

start




Partition


start




pivot




Partition




Partition




Partition




Partition




Partition



Partition



Partition



Partition

...

```
int* pbegin = begin, * pend = end;
```

```
while (true) {
```

```
    while (pbegin != pend && *pbegin < *pivot)
```

```
        ++pbegin;
```

```
    while (pbegin != pend && !(*--pend < *pivot));
```

```
    if (pbegin == pend) break;
```

```
    std::iter_swap(pbegin, pend);
```

```
    ++pbegin;
```

```
}
```

```
mid = pbegin;
```

...

Results



Sentinel Partition

- inner loops makes two checks per iteration:
 - determine if the end is reached
 - determine if the predicate applies
- can the check for the end be avoid?
 - use an object matching the predicate to stop

Sentinel on the Right

```
...
int* pbegin = begin, * pend = end;
while (true) {
    while (pbegin != pend && *pbegin < *pivot)
        ++pbegin;
    while (pbegin != pend && !(*--pend < *pivot));
    if (pbegin == pend) break;
    std::iter_swap(pbegin, pend);
    ++pbegin;
}
mid = pbegin;
...
```


Sentinel on the Right

```
...
int* pbegin = begin, * pend = end;
while (true) {
    while (*pbegin < *pivot)
        ++pbegin;
    while (pbegin != pend && !(*--pend < *pivot));
    if (pbegin == pend) break;
    std::iter_swap(pbegin, pend);
    ++pbegin;
}
mid = pbegin;
...
```

Initial Sentinel Partition

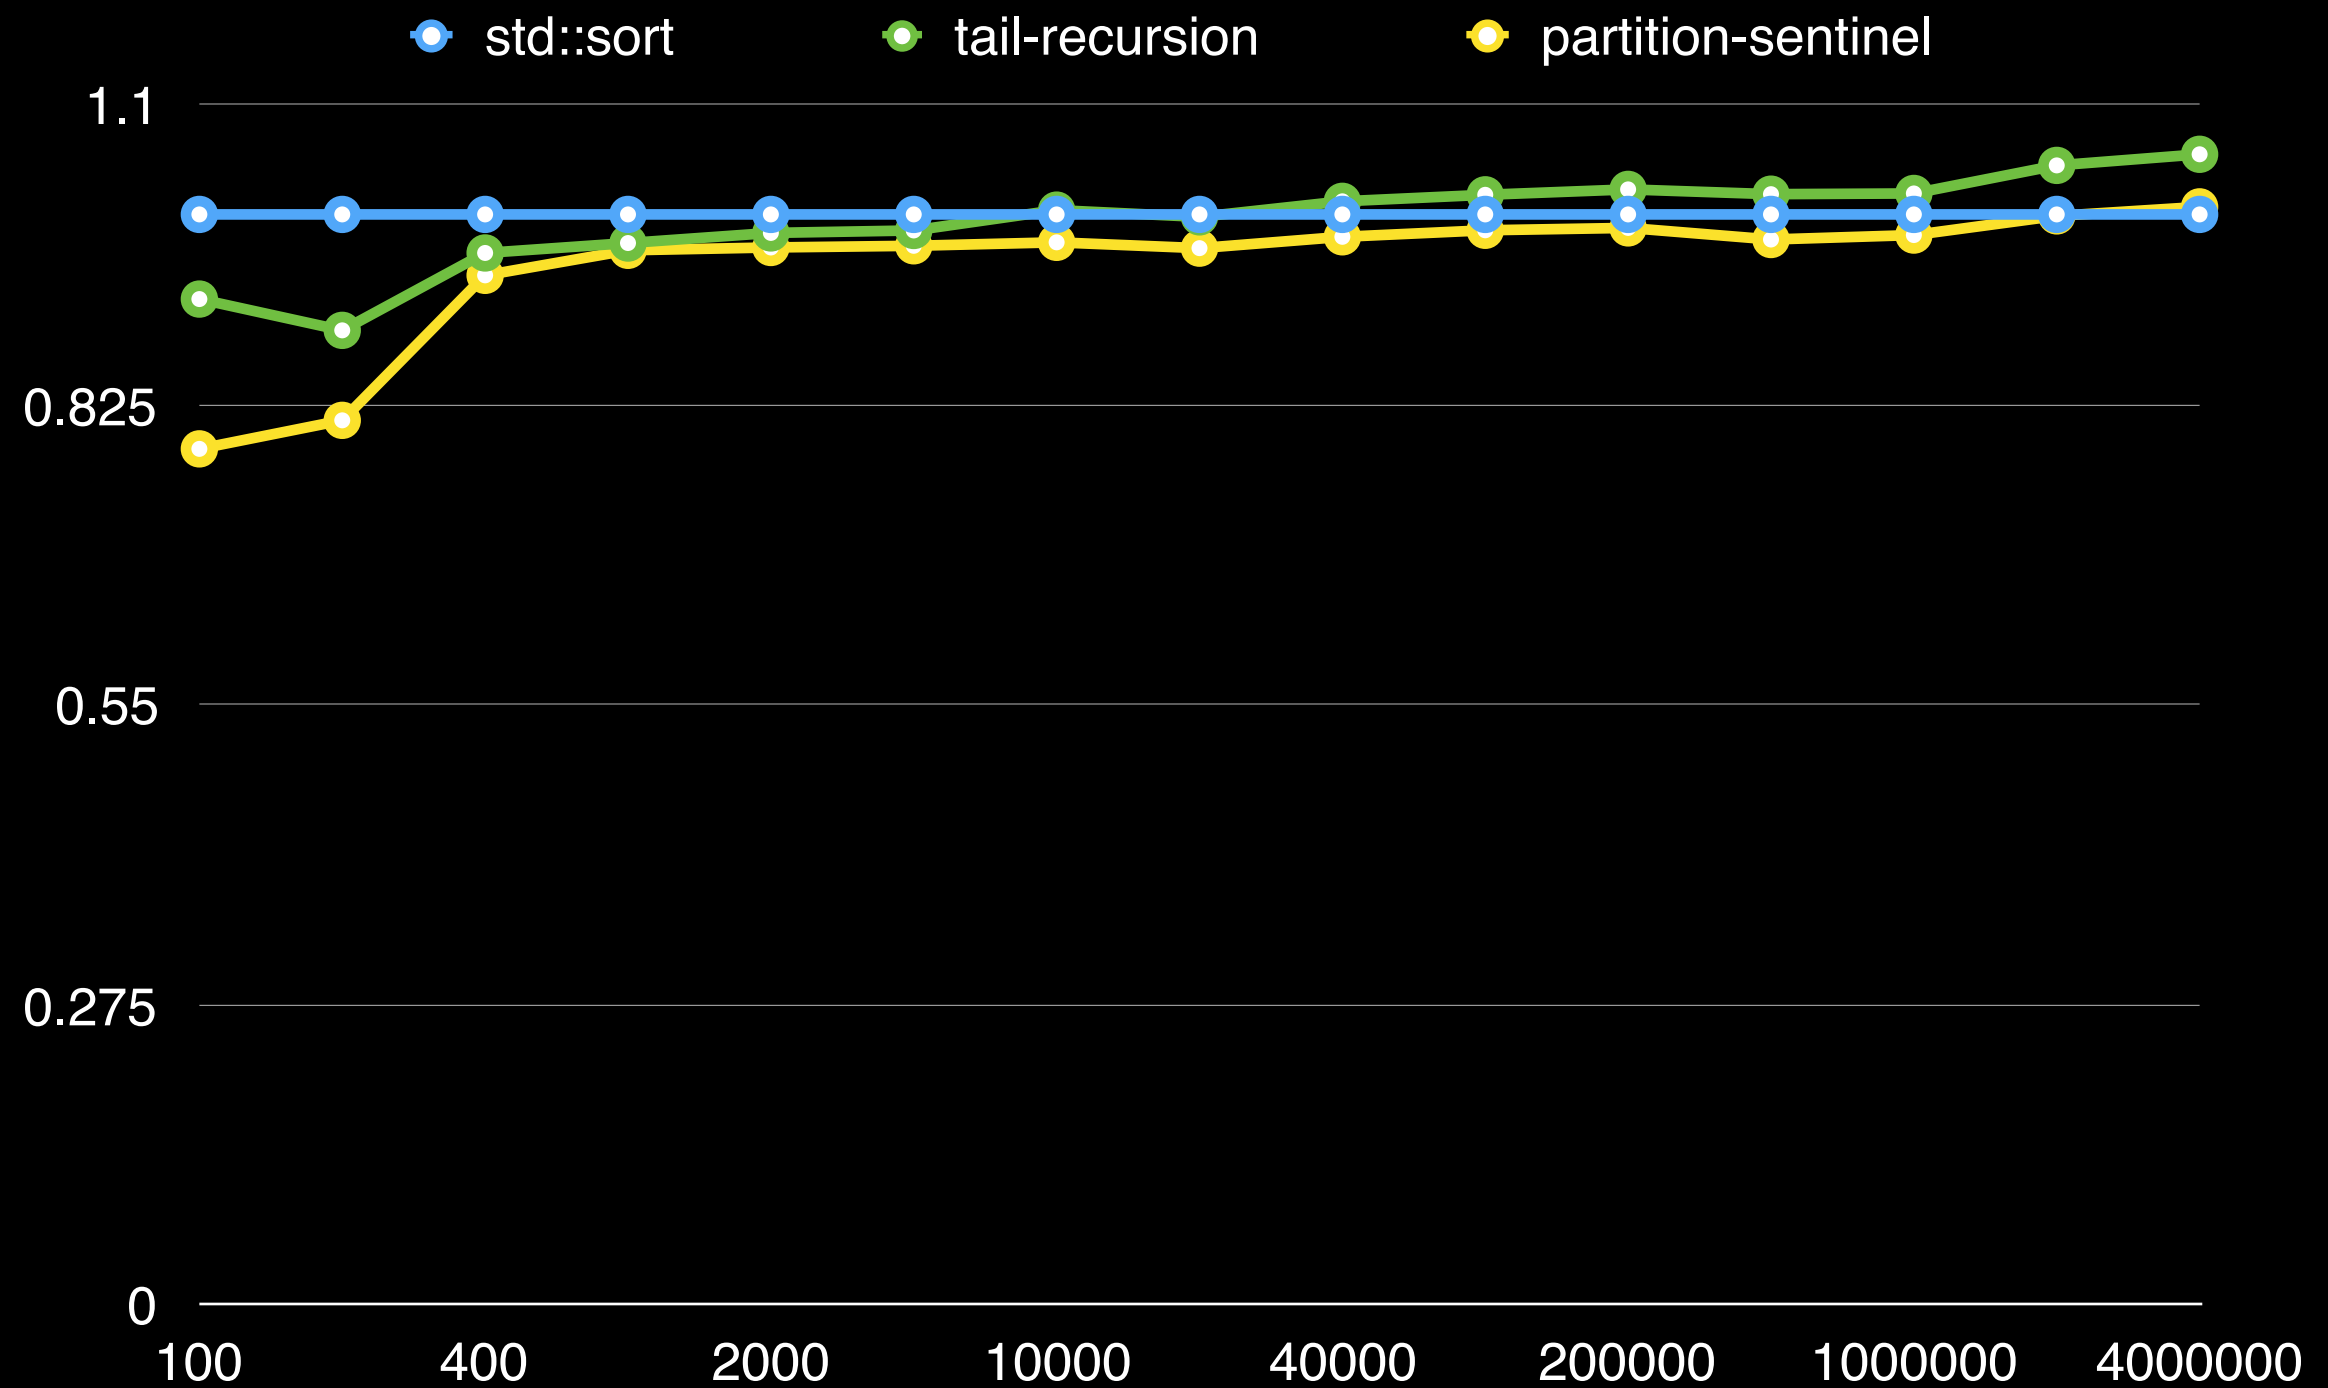
make sure there is a sentinel on the left side, too:

```
int* pbegin = begin, * pend = end;
if (pbegin != pend && !(*pbegin < p)) {
    while (pbegin != pend && !(*--pend < p)) {
    }
    if (pbegin != pend) {
        std::iter_swap(pbegin, pend);
        ++pbegin;
    }
}
```

Sentinel Partition Core

```
if (pbegin != pend) {  
    while (true) {  
        while (*pbegin < p) ++pbegin;  
        while (!(*--pend < p));  
        if (pend <= pbegin) break;  
        std::iter_swap(pbegin, pend);  
        ++pbegin;  
    }  
}  
mid = pbegin;
```

Results



Small Ranges

- many of the final ranges are small
- not worth to kick off an insertion sort
- manually sort ranges with up to 4 elements
- bigger ranges still use normal insertion sort


Small Range Dispatch

```
switch (std::distance(b, e)) {  
  case 0:  
  case 1: return;  
  case 2: if (*--e < *b) swap(*b, *e); return;  
  case 3: small::sort(b, b + 1, b + 2); return;  
  case 4: small::sort(b, b + 1, b + 2, b + 3); return;  
}
```

Sorting Four Elements

```
void small::sort(l p0, l p1, l p2, l p3) {  
    sort(p0, p1, p2);  
    if (*p3 < *p2) {  
        if (*p3 < *p1) {  
            auto tmp = *p3; *p3 = *p2; *p2 = *p1;  
            if (tmp < *p0) { *p1 = *p0; *p0 = tmp; }  
            else *p1 = tmp;  
        }  
        else swap(p2, p3);  
    }  
}
```

Results



Partition Did Nothing

- the partition may not do anything
 - sequence is sorted already
 - notably: all elements are the same
- gamble time on testing if it is sorted
 - ... and fix a few local inconsistency

Monitor Swaps


splitting the check may trigger small range optimisation

```
unsigned swaps = 0;  
... // partition incrementing swaps  
  
if (swaps == 0  
    && incomplete_insertion(begin, mid)  
    && incomplete_insertion(mid + 1, end)) {  
    return;  
}
```

Partial Insertion Sort

```
switch (std::distance(b, e)) { ... }  
for (int* it = b, cout = 0; ++it != e; )  
    if (*it < *(it - 1)) {  
        if (++count == limit) return false;  
        int* c = it, t = *c; *c = *(c - 1);  
        while (b != --c && t < *(c - 1)) {  
            *c = *(c - 1);  
            if (++count == limit) { *(c - 1) = t; return false; }  
        }  
        *c = t;  
    }  
return true;
```

Results



Putting It All Together

```
bool incomplete_insertion(int* begin, int* end) {
    switch (std::distance(begin, end)) {
    case 0:
    case 1:
        return true;
    case 2:
        if (*--end < *begin) std::swap(*begin, *end);
        return true;
    case 3:
        small::sort(begin, begin + 1, begin + 2);
        return true;
    case 4:
        small::sort(begin, begin + 1, begin + 2, begin + 3);
        return true;
    case 5:
        small::sort(begin, begin + 1, begin + 2, begin + 3, begin
+ 4);
        return true;
    }
    constexpr int limit = 8;
    int count(0);
    for (int* it = begin; ++it != end; )
        if (*it < *(it - 1)) {
            if (++count == limit) return false;
            int* c = it, t = *c;
            *c = *(c - 1);
            while (begin != --c && t < *(c - 1)) {
                *c = *(c - 1);
                if (++count == limit) {
                    *(c - 1) = t;
                    return false;
                }
            }
            *c = t;
        }
    return true;
}

void sort(int* begin, int* end, int depth, int max) {
    std::size_t len;
    while (20 < (len = std::distance(begin, end))) {
        if (++depth < max) {
            int* pivot = end - 1, * mid = begin + len / 2;
```

```
            std::swap(*mid, *pivot);

            unsigned swaps = 0;
            int p = *pivot;
            auto pred = [](int arg){ return arg < p; };
            int* pbegin = begin, * pend = end;
            if (pbegin != pend && !pred(*pbegin)) {
                while (pbegin != pend && !pred(*--pend)) {
                }
                if (pbegin != pend) {
                    std::iter_swap(pbegin, pend);
                    ++swaps;
                    ++pbegin;
                }
            }
            if (pbegin != pend) {
                while (true) {
                    while (pred(*pbegin)) {
                        ++pbegin;
                    }
                    while (!pred(*--pend)) {
                    }
                    if (pend <= pbegin) {
                        break;
                    }
                    std::iter_swap(pbegin, pend);
                    ++pbegin;
                    ++swaps;
                }
            }
            mid = pbegin;
            std::swap(*mid, *pivot);

            if (swaps == 0
                && incomplete_insertion(begin, mid)
                && incomplete_insertion(mid + 1, end)) {
                return;
            }
            if (mid - begin < end - mid) {
                sort(begin, mid, depth, max);
                begin = mid + 1;
```

```
            }
            else {
                sort(mid + 1, end, depth, max);
                end = mid;
            }
        }
        else {
            std::stable_sort(begin, end);
            return;
        }
    }
    switch (std::distance(begin, end)) {
    case 0:
    case 1:
        return;
    case 2:
        if (*--end < *begin) std::swap(*begin, *end);
        return;
    case 3:
        small::sort(begin, begin + 1, begin + 2);
        return;
    case 4:
        small::sort(begin, begin + 1, begin + 2, begin + 3);
        return;
    }
    for (int* it = begin; ++it != end; )
        if (*it < *(it - 1)) {
            int* c = it, t = *c;
            *c = *(c - 1);
            while (begin != --c && t < *(c - 1))
                *c = *(c - 1);
            *c = t;
        }
    return;
}

void sort(int* begin, int* end) {
    std::size_t size(std::distance(begin, end)), max(0);
    while (size >>= 1) { ++max; }
    sort(begin, end, 0, 2 * max);
}
```

Putting It All Together

```
bool incomplete_insertion(int* begin, int* end) {
    switch (std::distance(begin, end)) {
    case 0:
    case 1:
        return true;
    case 2:
        if (*--end < *begin) std::swap(*begin, *end);
        return true;
    case 3:
        small::sort(begin, begin + 1, begin + 2);
        return true;
    case 4:
        small::sort(begin, begin + 1, begin + 2, begin + 3);
        return true;
    case 5:
        small::sort(begin, begin + 1, begin + 2, begin + 3, begin
+ 4);
        return true;
    }
    constexpr int limit = 8;
    int count(0);
    for (int* it = begin; ++it != end; )
        if (*it < *(it - 1)) {
            if (++count == limit) return false;
            int* c = it, t = *c;
            *c = *(c - 1);
            while (begin != --c && t < *(c - 1)) {
                *c = *(c - 1);
                if (++count == limit) {
                    *(c - 1) = t;
                    return false;
                }
            }
            *c = t;
        }
    return true;
}

void sort(int* begin, int* end, int depth, int max) {
    std::size_t len;
    while (20 < (len = std::distance(begin, end))) {
        if (++depth < max) {
            int* pivot = end - 1, * mid = begin + len / 2;
```

```
            std::swap(*mid, *pivot);

            unsigned swaps = 0;
            int p = *pivot;
            auto pred = [=](int arg){ return arg < p; };
            int* pbegin = begin, * pend = end;
            if (pbegin != pend && !pred(*pbegin)) {
                while (pbegin != pend && !pred(*--pend)) {
                }
                if (pbegin != pend) {
                    std::iter_swap(pbegin, pend);
                    ++swaps;
                    ++pbegin;
                }
            }
            if (pbegin != pend) {
                while (true) {
                    while (pred(*pbegin)) {
                        ++pbegin;
                    }
                    while (!pred(*--pend)) {
                    }
                    if (pend <= pbegin) {
                        break;
                    }
                    std::iter_swap(pbegin, pend);
                    ++pbegin;
                    ++swaps;
                }
            }
            mid = pbegin;
            std::swap(*mid, *pivot);

            if (swaps == 0
                && incomplete_insertion(begin, mid)
                && incomplete_insertion(mid + 1, end)) {
                return;
            }
            if (mid - begin < end - mid) {
                sort(begin, mid, depth, max);
                begin = mid + 1;
```

```
            }
            else {
                sort(mid + 1, end, depth, max);
                end = mid;
            }
        }
        else {
            std::stable_sort(begin, end);
            return;
        }
    }
    switch (std::distance(begin, end)) {
    case 0:
    case 1:
        return;
    case 2:
        if (*--end < *begin) std::swap(*begin, *end);
        return;
    case 3:
        small::sort(begin, begin + 1, begin + 2);
        return;
    case 4:
        small::sort(begin, begin + 1, begin + 2, begin + 3);
        return;
    }
    for (int* it = begin; ++it != end; )
        if (*it < *(it - 1)) {
            int* c = it, t = *c;
            *c = *(c - 1);
            while (begin != --c && t < *(c - 1))
                *c = *(c - 1);
            *c = t;
        }
    return;
}

void sort(int* begin, int* end) {
    std::size_t size(std::distance(begin, end)), max(0);
    while (size >>= 1) { ++max; }
    sort(begin, end, 0, 2 * max);
}
```

Simple Implementation

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1;  
    int* mid = std::partition(begin, pivot,  
                             [=](int value){ return value < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```

More Generic

```
void sort(int* begin, int* end) {  
    std::size_t len = std::distance(begin, end);  
    if (len <= 1) return;  
    int* pivot = end - 1;  
    int* mid = std::partition(begin, pivot,  
                             [=](int value){ return value < *pivot; });  
    swap(*mid, *pivot);  
    sort(begin, mid);  
    sort(mid + 1, end);  
}
```


More Generic

```
template <typename T>
void sort(T * begin, T * end) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    T * pivot = end - 1;
    T * mid = std::partition(begin, pivot,
                             [=](T value){ return value < *pivot; });
    swap(*mid, *pivot);
    sort(begin, mid);
    sort(mid + 1, end);
}
```

More Generic

```
template <typename Iter>
void sort(Iter begin, Iter end) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    Iter pivot = end - 1;
    Iter mid = std::partition(begin, pivot,
        [=](auto&& value){ return value < *pivot; });
    swap(*mid, *pivot);
    sort(begin, mid);
    sort(mid + 1, end);
}
```

More Generic

```
template <typename Iter>
void sort(Iter begin, Iter end) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    Iter pivot = end - 1;
    Iter mid = std::partition(begin, pivot,
        [=](auto&& value){ return value < *pivot; });
    swap(*mid, *pivot);
    sort(begin, mid);
    sort(mid + 1, end);
}
```

More Generic

```
template <typename Iter, typename Comp>
void sort(Iter begin, Iter end, Comp c) {
    std::size_t len = std::distance(begin, end);
    if (len <= 1) return;
    Iter pivot = end - 1;
    Iter mid = std::partition(begin, pivot,
        [=](auto&& value){ return c(value, *pivot); });
    swap(*mid, *pivot);
    sort(begin, mid, c);
    sort(mid + 1, end, c);
}
```

Summary

- quick sort is fast
 - when pulling a lot of tricks
 - for the expected cases
 - it is a hybrid using multiple algorithms
- => real-world algorithms need a bit of tuning

Questions?